Design and Commissioning of the AARTFAAC all-sky monitor

Peeyush Prasad^{1,2}, Folkert Huizinga², John Romein¹, Daniel van der Schuur¹, and Ralph Wijers²

- ¹ Universiteit van Amsterdam
- ² ASTRON, The Netherlands Foundation for Radio Astronomy

Received <date> / Accepted <date>

ABSTRACT

The Amsterdam-ASTRON Radio Transients Facility And Analysis Center (AARTFAAC) array is a sensitive, all-sky radio imager based on the Low Frequency Array (LOFAR). It generates images of the low frequency radio sky in near real-time with spatial resolution of 10s of arcmin, MHz bandwidths and a time cadence of a few seconds. The image timeseries are then monitored for short and bright radio transients. On detection of a transient, low latency triggers will be generated for LOFAR, which can carry out follow-up observations. In this paper, we describe the implementation of the instrumentation, and its capabilities.

Key words. Radio Interferometry - Imaging - Radio Transients - Correlators

1. Introduction

- What are celestial transients

Transient astronomy deals with the detection and characterization of celestial transients, sources in the sky whose detectable properties can change on short timescales. These explosive events provide insight into a variety of astrophysics, ranging from emission mechanisms of jets to properties of the intervening medium [ref]. There is a rich history of the detection of such objects across wavelength ranges, with each wavelength regime probing a different parameter space [ref].

- Current state of radio instrumentation

Studies of short duration radio emission from such objects has been restricted to either very short timescales (milliseconds to seconds, e.g. Pulsars), or to comparitively longer timescales (months to years) primarily due to instrumental or observational time constraints, the latter due to the narrow fields of view. Radio instrumentation is available in broadly two classes; single dish or phased array beam formed timeseries characterized by high time and frequency resolution, fields of view of a few degrees but poor spatial resolution. Aperture synthesis imaging observations address the latter to provide high spatial resolution, but have poor time resolution, typically needing several hours of observation time to build up adequate coverage in the UV plane via earth rotation aperture synthesis. The above classes roughly translate to coherent (short timescales) and incoherent (longer timescales) sources of emission.

- Current state of the science

The serendipitous discovery of a new class of radio transient termed Fast Radio Bursts (FRBs) has galvanized interest in the field. The detected FRBs are characterized by high associated dispersion measures, high brightness and short timescales. They are non-repeating for the most part. Their unknown origin requires not only their discovery, but also rapid followup over a large wavelength regime to establish emission phenomena and associated parameters.

The last requirement has led to the development of large field of view radio sky monitors, with an aim of continuously sur-

- veying large parts of the visible sky with shallow sensitivity and at high time resolution. A trigger is generated on the reliable detection of a transient in close to real-time, allowing other telescopes to carry out follow-up observations.
- Suitability of low frequency observations
 TODO: Paragraph about transient sources, some stuff about spectral indices of coherent emission and expected class of sources, at what brightness levels can we expect to see things (take from Lazio LWA paper).>
 - <TODO: Paragraph summarizing the current state of knowledge of low frequency transients: Stewart NCP transient, other searches at low freq. Conclusion: Need for more monitoring.>
- What is the AARTFAAC

The AARTFAAC radio transient monitor is an All-sky radio telescope based on LOFAR. Its goal is to continuously scan the skies for bright transients, and on reliably detecting one, to generate a trigger to other telescopes for sensitive, broad band monitoring. It is a leading effort among a group of new radio telescopes aiming for detection of bursts of radio emission by continuous monitoring of the low radio frequency sky. Such telescopes are characterized by having moderate resolution and sensitivity as compared to contemporary telescopes, but with extremely wide fields of view (typically all sky), high availabilities and autonomous calibration and imaging in near real-time.

The latter requirements are unconventional by contemporary Radio Astronomy standards, and make their implementations challenging. The antenna elements used to achieve the wide fields of view are typically dipoles, however, their low individual sensitivities requires an order of magnitude larger number of elements in the array. Bringing the resulting large number of data streams to a central location, as well as their correlation for carrying out aperture synthesis imaging in real-time thus poses a significant I/O and compute challenge. Further, the wide fields of view at the sensitivities of operation also result in direction dependent effects on the incoming signals, mostly due to the ionosphere. These pose a

challenge to calibration, especially when carried out in an autonomous manner.

Apart from its primary goal of trigger generation on the detection of transients, the telescope products find use in a variety of science cases. These include wide field ionospheric monitoring via apparent flux and position variations of calibrator sources, Solar monitoring, RFI surveying, LOFAR beam model validation etc.

AARTFAAC as a data transport, reorganization and computing problem.

The wide field of views necessary for an instrument like AARTFAAC can be achieved by sampling the sky with wide field dipoles. This, however comes at the cost of lowered sensitivity per receiving element. An instantaneously well sampled UV plane is needed to generate a PSF with low sidelobes. Both requirements can be met by spatially spreading a large number of dipoles. The highest sensitivities can also be achieved by the coherent correlation of the incoming signal, requiring access to the nyquist sampled signal at full resolution. Such an arrangement then requires the aggregation of high bandwidth data from the receiver elements, necessitating a high speed data network.

The incoming sampled voltages pass through various signal processing blocks, resulting in the generation of light curves for sources in the image. The estimation of spatial coherences requires the reordering of data to make optimum usage of compute resources. Thus, the functioning of the telescope depends on the optimization of the data transport, data reorganization and computing using available resources. An advantage of having an operating model consisting of signal processing blocks operating on high resolution data is the configurability of the telescope into different observing modes, as well as the tapping off of data from an upstream location. The latter ability makes a piggy-back instrument like the AARTFAAC possible. An important resource to optimize is the development time for each data routing or processing block, and this has been taken into consideration in the AARTFAAC.

Paper description

In this paper, we describe the AARTFAAC telescope system architecture, its instrumentation, and the commissioning of its various subsystems. Section 2 describes the array and the receiving antenna elements, its relationship with LOFAR, and introduces the full architecture of the instrument. Section 3 describes the hardware implementation in the field which allows creating a data path in parallel to LOFAR. This makes AARTFAAC processing independent of LOFAR to a large extent. In Section 4, we describe the implementation of a real-time, GPU based correlator for AARTFAAC, while Section 5 details the real-time, autonomous calibration and imaging implementation. Section 6 describes our control system for the full instrument, which also interfaces with LOFAR. In Section 7 we present performance metrics of the instrument as a whole.

2. The AARTFAAC array

We begin by summarizing the subsystems of the LOFAR telescope relevant for AARTFAAC processing in Section 2.1, and then elaborating on the scheme for creating a coupled data path for independent processing by AARTFAAC.

2.1. LOFAR telescope architecture

The LOFAR telescope ?? is a new generation radio interferometer covering the frequency range from 10-90 MHz using inverted V-dipoles known as Low Band Antenna (LBA), and from 110-240 MHz using Bowtie dipoles, also known as High Band Antenna (HBA). The antenna are linearly polarized, being made up of orthogonally placed dipoles in the E and H plane. The LBA dipole has a sensitivity pattern with a 6dB field of view of about 120°, while the HBA dipoles first undergo an analog phasing within a 4x4 tile, which results in a field of view of about TODO. Due to this restriction, the AARTFAAC array utilizes only the LBA part of the telescope.

The telescope itself consists of a large collection of antennas, spatially organized into several 'stations', each spread over 60m. The stations are laid out in a dense core: 24 (TODO: Check) stations within a 2km radius, while the long baselines of stations of upto a 1000km are also present. At the station level, the received and conditioned analog signals from a dipole are baseband sampled with a 200MHz clock and 10-bit quantized.(TODO: Check). The signal from each polarization is then split into spectral subbands of 200kHz via a polyphase filterbank implementation. In the regular LOFAR station level processing, the dipole subbands are then digitally phased in hardware towards the direction of an astronomical source to form a station beam, which is then transmitted over optical fiber for further interferometric processing with other stations.

The LOFAR correlator architecture needs real-time processing to reduce the large volume of data being produced. Its implementation is also based on a hybrid architecture using GPG-PUs and CPUs. [TODO: Should this be mentioned at all? Might weaken the case, with our implementation being very similar to the LOFAR case.]

A schematic representation of the LOFAR level processing is shown in Figure. TODO

2.2. The AARTFAAC system

The AARTFAAC array consists of 12-stations from within the core of the LOFAR telescope, with interdipole distances ranging from (TODO) within a station, and a maximum of TODO across stations. Due to the requirement of dipole level data in order to achieve all-sky imaging, the AARTFAAC creates a coupled data path to an independent processing architecture, prior to the phasing up of the dipoles in the LOFAR processing flow. This allows simultaneous observing with LOFAR, leading to high availability of the AARTFAAC system. A subset of the available subbands are correlated in a dedicated GPU based correlator in realtime. The estimated visibilities are sent over TCP/IP to dedicated servers for carrying out the autonomous and real-time calibration and imaging. The generated images are further sent to a software pipeline for the actual detection of transients, based on comparison of the image timeseries. A (planned) trigger generation subsystem will publish reliable triggers in the form of VOEvents [refTODO], which can be claimed by other telescopes to observe candidates with high sensitivity and resolution.

2.2.1. Array configuration

The choice of stations is dictated primarily by imaging quality and sensitivity, but also due to constraints on the latency of calibration and imaging. The central six stations of the LOFAR telescope (called the superterp) form a densely sampled UV plane, and are ideal for wide field imaging due to their being co-

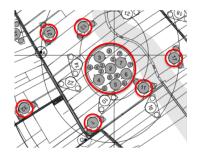


Fig. 1: The spatial distribution of AARTFAAC-12 stations within the core of LOFAR stations.

planar to high accuracy (centimeter level). The outer six stations provide higher sensitivity and resolution. The salient features of the LBA_OUTER station configuration for the chosen stations are shown in table TODO. Figure 1 shows the LOFAR stations that are part of the AARTFAAC system.

TODO: Add 12-station beam characteristics, expected confusion noise contribution.

The station consititute the first component of the radio sky monitor, and are the only components shared with LOFAR. The AARTFAAC monitor consists of further subsystems which are independent of LOFAR processing. Its overall architecture is shown schematically in Figure 2, and illustrates the main processing subblocks of the instrument. To summarize, a user selectable subset of subbands from every dipole is transferred as UDP packets over a dedicated 10Gbit fiber connection to the central processing systems. These are received by a streaming software correlator implementation which aligns the data and estimates the spatial covariance matrix between every pair of dipoles. The generated visibilities are streamed over TCP/IP to a calibration and imaging pipeline component which carries out autonomous imaging. The images are then analyzed by a software tool (The Transients Pipeline, TraP), which extracts the light curves of sources within the image, and analyses them for variability using a number of parameters. The TraP is described in more detail in ??.

The specifications of the AARTFAAC monitor are listed in Table TODO. We describe the various subsystems making up the AARTFAAC All-sky monitor in the following sections.

3. Remote Station Level Processing:

Remote Station processing refers to instrumentation installed in the field coupled to the receiving antennas post balun. This consists of Digital Converter Unit (DCU) boards and the Remote Station Processing (RSP) boards. The former caters to analog signal conditioning of the received voltage, followed by base band digitization at 200MHz, with TODO bit resolution.

The RSP boards handle the reception of sampled voltages from the DCU boards and their first stage processing. In the following, we describe the parts relevant to the functioning of the AARTFAAC telescope. The RSP board consists of TODO NUM TODO TYPE FPGAs, with TODO input bandwidth over TODO connector, and TODO output bandwidth over TODO connector. [TODO Add reference] describes the implementation of this board in more detail. A single RSP board can handle the processing of sampled data from 4 dual polarized dipoles. Since a LOFAR station is made up of 48 dual polarized dipole antennas, 12 such boards are required per station.

Signal processing: The RSP board carries out the first stage polyphase filter bank implementation common to the LOFAR and AARTFAAC telescopes. This filter bank analyzes the voltage timeseries sampled at 200 MHz available from the samplers, into a complex voltage spectrum of 512 subbands. Thus, the entire analog band of the LBA between 10-90MHz is available for further processing. The output, for a set of 1024 real voltage samples, consists of fixed point complex values per subband. These values fundamentally have a 16-bit resolution on the real and imaginary components. [Add some more on dynamic range expected etc.]

Station level beamforming in a particular direction requires calculating the weighted sum of all dipoles of the same polarization, with the applied weights being dependent on the direction of beamforming. Each RSP board fundamentally generates the beamformed product for its 4 client dipoles for every subband. These are termed as 'beamlets'. The necessary exchange of beamlets with other RSP boards to create the station beam is achieved via the interconnect between the various RSP boards.

Each data product from the RSP signal processing is encapsulated in a packet, and marked with a datatype magic number. This allows differentiation between the generated data products.

Interconnect: A ring network consisting of four 10-GigE links interconnects the RSP boards to each other, as shown in Fig. 4. It is implemented by daisy chaining the GigE links of the RSP boards to each other. It carries beamlet data, and has an AARTFAAC specific mode (enabled only on AARTFAAC boards? TODO check) where the raw subbands for every dipole polarization are also transferred onto the ring network. Of the total TODO bandwidth of the ring network, about TODO is occupied by LOFAR specific products. The remaining bandwidth carries the per dipole subbands, which are used exclusively for AARTFAAC processing.

The fundamental limitation to AARTFAAC processed bandwidth is presented by the interconnect, and depends on the bitmode chosen. The bandwidth available to AARTFAAC is limited to 36 subbands in 16-bit mode, 72 subbands in 8-bit mode, or 144 subbands in 4-bit mode. The bit modes are mutually exclusive, and can be set via control registers. A random selection of subbands from the available 512 can be inserted into the available slots on the interconnect. This selection can be made via manipulation of the control registers of the RSP board. This allows AARTFAAC to achieve high sensitivity by placing subbands contiguously, and later integrating them, while at the same time achieving spectral coverage by placing subbands to sample a larger extent of the analog spectrum.

Available bit modes: The system offers the ability to trade off dynamic range in the polyphase filter bank outputs with the

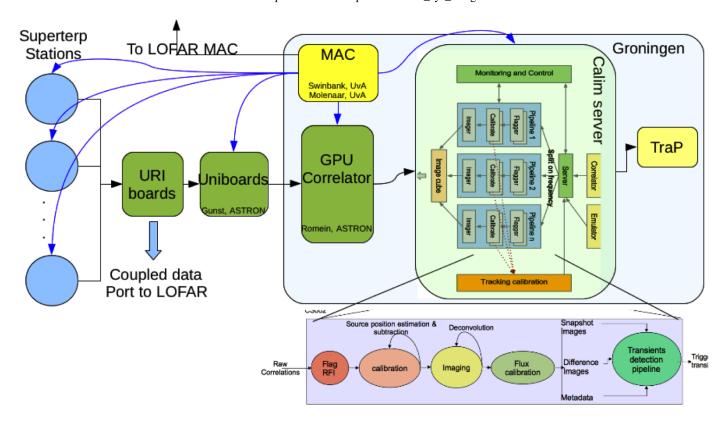


Fig. 2: Overall architecture of the AARTFAAC all-sky monitor depicting each processing subblock.

number of subbands available for further processing. This is done by reducing the number of allocated bits to the real and imaginary components of the subband outputs, leading to an increased number of subbands. The bit mode of AARTFAAC can be set completely independently of LOFAR's choice of bit mode. The choice between the various bit modes depends on the RFI environment of the observation.[TODO: How is the data converted from 16-bit to 8-bit? Which 8-bits are taken as the output? Implications?]. An 8-bit complex representation of the filterbank outputs is found to be adequate for almost all observing conditions except during severe RFI.

Sampling clock and Timing: A clock distributer board [TODO REF] is used to distribute a 10MHz reference to every one of the 12 AARTFAAC stations. This reference is generated by a GPS discplined Rubidium frequency standard. This ensures that an identical (hence coherent) clock is used for the sampling of data from the AARTFAAC stations.

Every station is equipped with a Local Control Unit (LCU), which is an embedded (TODO:Check?) processor running a Linux operating system. These systems are networked to the LOFAR control system, and also act as NTP clients. Thus, their absolute times are aligned to better than a few milliseconds. This absolute time from the station LCU is communicated to the RSP board [TODO: Check?] on station reset, as the reset value to a 64bit counter. Once set, the station hardware updates this counter on a derivate of the available 200MHz reference, thus ensuring that the absolute time is embedded in the data with few nanosecond [TODO check] resolution. All further aligning and timing of the incoming data is carried out based on this embedded timestamp.

The Uniboard-RSP Interface (URI) board: This AART-FAAC specific board creates a coupled path for the subband data

by segregating them from the LOFAR specific data products flowing on the station ring network. This is implemented by interfacing the interconnet of 4 RSP boards to a single URI board, which interrogates the incoming packets, and selects the subband outputs based on packet type. This block further implements the first stage of the necessary swizzle operation by statically routing a subband from all dipoles of the station to a single output lane.[Add more detail here.]

All together, three URI boards are adequate to transfer and swizzle 36 subbands at 16-bits into the uniboard based router.

The Uniboard based data router: This data routing unit is the interface between the station level instrumentation, and the next signal processing unit, the correlator. It fundamentally consists of 4 upstream (called backnodes) and 4 downstream (called frontnode) FPGAs. Each of the backnode FPGAs receives 8 consecutive subbands out of the 36 subbands in the URI board output, making 32 subbands available to the Uniboard. A second level of swizzling is carried out at this stage via physical routing of subbands between the URI board and the Uniboard. The resulting stream out of a backnode FPGA consists of 8 subbands of all 48 dipoles of a station. These data are transported to a single Frontnode FPGA. The latter encapsulates the data into a UDP packet which is transmitted on a long haul 10Gigabit Ethernet interface to the remote correlator. A single Uniboard operates two 10Gbps links to the central processing machines, located at the University of Groningen about 50Km away.

Subband data format: (Is this level of detail needed?) **Commissioning effort and Performance:**

Data rates and bandwidth limitations: The stations operate on a fixed sampling clock of 200 MHz, leading to an output rate of 12 Mbps per dual polarised dipole antenna per 16-bit complex subband. The limited ring network bandwidth of a sta-

tion allows only 36 of the available 512 subbands of 16-bits from all dipole antennas (total bandwidth 20Gbps) to be carried to the URI board. The station Uniboard further restricts the sampled bandwidth to match its output bandwidth into 2x10Gbps links. Of the incoming 36 subbands, only 16 are forwarded for correlation. Each 10Gbps link carries 8 subbands, corresponding to 4.5Gbps of data.

Control interface: The control of the remote station electronics consists of two layers. At the FPGA level, command and status registers have been opened up. These can be accessed via the a separate (TODO: Double check) control Gigabit Ethernet interface. Each station is also equipped with a Local Control Unit (LCU) computer which provides an abstraction layer to the hardware. All control and monitoring commands from a global control system are addressed to the LCU, which provides a tool that can access the hardware registers via a driver, which ultimately communicates the commands over the Gigabit ethernet control link to the RSP boards of the station.

Figure 4 depicts the station level ring network, whose bandwidth is shared between the beamformed subbands as well as the dipole level subbands. The ring network bandwidth constrains the processed AARTFAAC bandwidth to a fundamental maximum of 36 16-bit subbands, or about 7 MHz, while the Uniboard interface further restricts the bandwidth to 8 subbands per 10Gpbs output link. The URI boards in combination with the uniboards carry out a first level of the incoming data transposition.

4. The AARTFAAC real-time correlator

- What does the correlator do

The correlator subsystem estimates the spatial coherence between all pairs of the 1156 polarizations of the drooped dipoles, per subband. This is done by computing the average cross correlation between antenna polarization per subband. Its input is formed of the subbanded complex voltage timeseries from each polarization of every antenna. The output consists of a stream of visibilities with a chosen time and frequency averaging.

The correlator requirements for AARTFAAC The AARTFAAC correlator must estimate the correlation from 1152 input streams with as little latency as possible. It thus needs to ingest a very high data rate. Further, it should be able to produce channelized output due to requirements of RFI flagging, while being configurable in its time integration capabilities. A major requirement of the implementation was that of development time, which effectively eliminated a

- Motivating the GPGPU approach

firmware based approach.

Of the available approaches, an heterogenous GPGPU approach turned out to be the best match between our requirements of ease of development and performance. The correlation operation has a low Arithmetic Intensity, implying that the data brought to a device ALU is operated on only a few times. This, combined with the relatively restricted I/O bandwidth between the host CPU and the device GPU, has been used as an argument against using GPGPU systems for correlation. However, with increasing number of signal sources, the I/O requirements scale linearly, while the compute requirements scale quadratically. Thus, a large enough input system can make the application compute bound. This is the case withthe AARTFAAC system. Assuming the input subbands from the 1152 dipole elements would be split into 64

channels, with a one second time integration, the total required compute power required is 17 TFLOP/sec. The latest GPUs come close to this requirement.

Functional blocks

The signal processing blocks necessary to achieve the correlator specs are shown in Fig. TODO. After the incoming subbanded data from each station is aligned, a Polyphase filter bank is implemented in order to increase frequency resolution. The resulting channels are then delay compensated for the fixed cable delays in the system, as well as for the bandpass response of the first stage Polyphase filter bank. Following this, signals from corresponding timeslices for all stations are correlated. The generated covariance matrix is then written out for further processing.

4.1. Implementation Hardware architecture:

The AARTFAAC correlator has been implemented using a GPGPU approach, with a server class machine utilizing a GPU device for carrying out the computation necessary for the correlation. The host CPU based machines acts as the interface between the stations and the GPU devices. They implement the data reception, and arbitrate the data distribution between different GPUs. They also carry out the last stage of the swizzle operation to optmize the data layout for computing the spectral coherence on a GPU.

The implementation consists of two identical machine configurations, each capable of handling 16 incoming subbands from 12 stations. Each machine consists of dual Xeon-class processors with 24 cores each. The CPUs are connected to 32 GB of memory each. Five AMD FirePro S10000 GPU cards are available per machine, interfaced to the CPUs over 8x16bit PCIe lanes (TODO: Verify correctness). The memory is organised in a NUMA configuration, with a QuickPathInterface (QPI) connection between the two CPUs. To receive the 5Gbps data output from each of the 12 stations, the server is equipped with 2x40Gbps infiniband interfaces, each collecting the 30Gbps data from 6 stations. The 5Gbps single station bandwidth is aggregated by a switch with dual 40Gbps output port, which in turn connects to an induvidual correlator server.

NUMA domains: In order to optimize throughput, a single machines resources are organized into two NUMA domains. Each domain includes a set of processing cores, the memory associated with them. The hardware splits the network I/O such that a single infiniband card can be present in each of the two NUMA domains. However, the split of the 5 GPU cards cannot be made symmetrically between the two domains, leaving one domain with 3 GPUs, while the other domain as 2 GPUs.

4.2. Implementation of functional blocks:

In this section, we provide a description of the implementation of individual functional blocks on our target hardware.

CPU Data alignment and transpose:

The raw data input packets from individual stations need to be collated for an entire integration period in host memory, before we can ship data to the correlator. We allocate 4 processor cores for the execution of a thread pool for I/O handling. Each thread handles the data reception from one station. A large buffer for the incoming dat is also setup. Each of these I/O thread manages reception of data from the network, and record data into the large buffer. While the copy of contiguous smilim work on it during ayourwhileFinally, we pay thilarge buffer space, correcting.

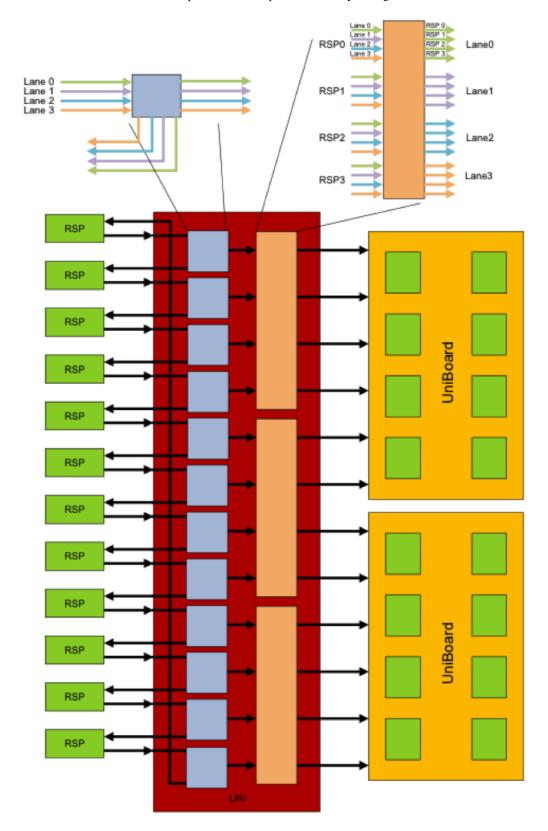


Fig. 3: The station level hardware changes allowing creation of a coupled data path for AARTFAAC data flow.

The incoming UDP packets are organized in the Station, Time, Subband order, whilte it is more efficient to rate the visibilities The visibilities.order.

The data buffer consists of a large number of integration units (TODO: Is this true?I think a command line parameter allows

specifying this, but then how is latency maintained?) The assembled data buffer is asynchronously transferred to the GPU device memory for further processing.(TODO: When is the last stage swizzle carried out? while transferring to device memory,

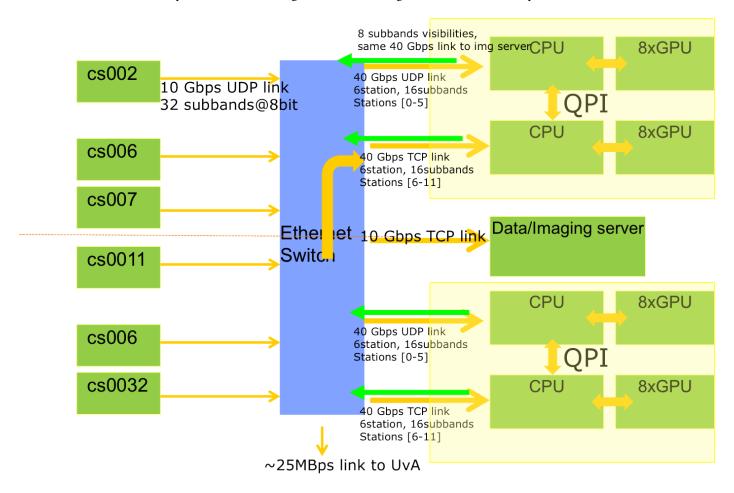


Fig. 4: The GPU correlator implementation using a pair of Xeon class server machines hosting AMD GPUs

or while writing the received data into host memory? Is it even relevant without first doing the PFB?)

Polyphase Filterbank kernel: A GPU thread is created to handle the processing of each subband separately. The first signal processing block is a PolyPhase filterbank, applied onto the single subband. A 64-tap FIR filter first increases the frequency resolution to 64 channels across the subband. This is followed by a 1-D complex to complex 64-point FFT on the channel data, leading to a final output frequency resolution of about 3kHz. Both the first stage FIR filter and the subsequent FFT is carried out on the GPUs. The FIR filter implementation is custom code, while the FFT is carried out using a stock (TODO: which one?) openCL library. TODO: Memory layout.

Delay and Bandpass compensation kernel: Subsequent to the PolyPhase filter bank implementation, a fixed delay compensation is applied to account for the cable delays of the dipoles with respect to a reference antenna. These delays are obtained via a separate calibration, which is typically carried out at a cadence of a few months. The delays are available in calibration tables, and the frequency resolution is high enough to apply them as phase rotations of the visibilities. The first stage polyphase filterbank implementation in the RSP board results in a deterministic amplitude modulation on the subband bandpass, and leading to unequal powers in each subband channel. This is demodulated via the application of a fixed amplitude correction by applying channel dependent weights. TODO: Memory layout.

Correlation kernel: Each dipole's subbanded and channelized data is then ready for correlation. For the dual-pol input corresponding to N input stream, the correlator produces a

covariance matrix of size 2Nx2N cells, with each cell corresponding to a product of the two polarizations, one of the XX, XY, YX and YY combinations possible.

Due to the low Arithmetic intensity of the correlation process, and the limited bandwidth between host and device memory, the GPU correlator implementation optimises for data locality, based on the available number of registers and per compute unit (TODO: technical term here?) [Ref: van Nieuwpoort Romein]. The covariance matrix is split into 2x2 cells, which allow a vectorised load of 4 floats (float4), with the accumulators remaining in registers. The computed correlations are then collated into device global memory by every thread group, with a final write from device to host memory. BIG TODO: Memory layout.

Asynchronous host to device transfers overlapping with compute: A source of throughput bottleneck for the streaming, low Aritmetic Intensity correlator application on a GPGPU platform is the bandwidth requirement for the transfer of large volume dipole data from host to device, and to a smaller extent, the transfer of computed correlations from the device memory to the host. The PCIe bandwidth limitation (TODO: Is it just bandwidth, or multiple access to the same memory address?) can be addressed by scheduling data transfer asynchronously, while overlapping the computation of the kernels with the data transfer.

Parallization axis: Parallelization is fundamentally on the frequency axis, with each subband being processed independently. Within a subband, each station data is processed independently till the correlation stage, where a synchronization barrier

is applied to time align all the dipole streams. Each channel of the subband correlation matrix is also processed independently via a GPU thread group. Further, a single channel covariance matrix is broken into cells, which are also computed independently.

- Enormous I/O requirements. Is a streaming, real-time application with large - I/O requirements. This can be a problem for implementation in many core - hardware. - Computational demands grow quadratically. Challenging computing step, since computation grows quadratically with number of inputs. - Recent trend is to correlate in software instead of hardware, due to flexibility and - to reduce development effort. - I/O requirements: 1156*8*16*195312.5 bits/sec. - Output resolution: Limited by time and frequency smearing effects, described more in a forward section. - Justification for correlating on a GPU as against special purpose hardware, or super computers (take from a GPU paper). - Logical blocks of the correlator: The correlator consists of the following logical blocks: - The input section: Job is to receive the incoming data. - The ring buffer for alignment of input streams: Forms a time aligned set of input data. All incoming data packets are copied into a slot in memory based on the timestamp from the data. - A second level polyphase filterbank is applied to the incoming subbanded data. This provides the necessary input frequency resolution, and also allows to carrout the subband amplitude modulation caused as a side effect of the first stage poly-phase filter bank. The correction is based on a theoretical response to the PFB. - The input is then passed on to the GPU via a host-device transfer, to carry out the actual correlation. The host-device transfer is on the critical path from latency and throughput perspective. - The combination of two stations is called a baseline, total number of baselines is N*(N+1)/2. This includes autocorrelations. - The correlator can operate on different input data modes. We describe the results in terms of 16bit complex inputs. We address the question of scaling up the system in section [forward ref]. - Description of the operation of correlation. - Description of the output products. - Choice of doing floating point operations (get info from e.g. D'Addario) -How many FLOPS/byte of incoming data? (arithmetic intensity). - Memory organization in host memory. - Memory organization in device memory. Any optimizations for reduction of memory loads. - Tuning of tile size to implementation architecture. Make the tilesize as large as possible while fitting in registerspace for maximum utilization of loaded data. - Description of tile selection of data and resulting arithmetic intensity. - Table of properties of the chosen architecture (like table 3 of van niewpoort and romein, many core correlator architecture.) - Ratio between FLOPS and Bytes/sec of memory bandwidth: indicator of performance of memory system. - Performance is bound both by theoretical peak performance, and the product of memory bandwidth and arithmetic intensity. - Mention theoretical peak performance, get teh actual value from John Romein (See Sec. 3.6 of Nieuwpoort paper), what fraction of peak performance is achieved? -Kernel performance? - Achieving GPU theoretical performance depends on setting up an adequate tile size, this in turn depends on the nnumber of available registers, and the host-device band-

Code related summary points – - Overall program architecture: What does the CPU do, what gets offloaded to the GPU?
What tile size is used? - Are the XY and YX hands calculated?
What does -m9 do? - Buffer sizes, memory footprint?

5. Real-time calibration and imaging

- Architecture, implementation choices, performance

- Visibility and image buffering strategy for followup analysis of detected transients
- Quicklook images data path
- Unit test architecture
- Interface to TraP

6. The AARTFAAC control interface

Figure 5 shows the functional blocks of the AARTFAAC control system, and their interface to the LOFAR scheduling system.

- Control system description
- Interface with LOFAR
- Monitoring interface: AARTFAAC webpage

7. Commissioning results

- Overall latency of the system
- Long term performance of the entire system based on logs.
- Performance in various bit-modes, with different number of subbands, expected sensitivity.
- Imaging quality Vs. latency: 6 station to 12 station.

8. Discussion

- Long term operations of the instrument.
- Triggering mechanism: VOEvents?
- False alarms (?): Known sources of transients

8.1. AARTFAAC Scalability

The AARTFAAC All-Sky monitor can be scaled up along several dimensions: The number of dipoles (spatially), the processed bandwidth (spectrally), and by increasing the imaging cadence (temporally). Spatiall scalability is

8.2. Impact of sharing dipoles with LOFAR on AARTFAAC

LOFAR operates using either the LBA or the HBA antenna at a time. Further, due to limited station level electronics for stations within the core, only a subset of the available station dipoles can be utilized. This implies that the AARTFAAC telescope is dependent on LOFAR for the choice of antenna and station configuration, reducing the availability for all-sky monitoring. Within the station, only the LBA_OUTER station configuration is currently deemded suitable for real-time imaging. This mode of LOFAR operation favorable to AARTFAAC depends on the observing schedule and the proposed observations. Table TODO shows some statistics from previous cycles on the fraction of observing modes favorable to AARTFAAC. Based on this, it may be resonable to expect AARTFAAC to operate TODO fraction of time, typically.

8.3. Hierarchical swizzling of incoming data:

The voltages are available as a sampled timeseries per polarization, with time samples laid out contiguously in memory. For efficient computing, these need to be transposed such that the same timeslice from all dipoles lie contiguously in memory. This operation involves transposing a matrix of data with a size NxN, with N being the number of data sources. This memory becomes inflated by a factor corresponding to the number of timesamples in the integration unit, as well as the number of

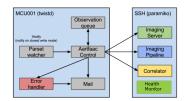


Fig. 5: The control system architecture which interfaces with the LOFAR observation scheduling system and triggers AARTFAAC observations.

polarizations being processed. Since the memory footprint can be quite large, this is quite an inefficient operation due to memory read latencies etc.

We achieve this by spreading the transpose (or swizzle) operation over the nodes on our network. This is done in the following manner:

The URI boards described in Sec. [TODO] carry out a first level of transposition by physically routing a subbands from a group of 4 dipoles to a single output. The second stage swizzle is carried out by the Uniboard. This is at the station level, and routes all dipole data from one station as a contiguously laid out group. Finally, the inter-station swizzling is carried out by the CPU host component of the correlator. Here, the PC memory is used to rearrange the incoming data and to carry out the transform. The resulting output is then optimally laid out for correlation over all dipoles of the AARTFAAC, for every polarization, and over the chosen integration time range.

8.4. Dynamice range and precision of computing

The 12-bit sampling (TODO: Confirm) provides for a dynamic range of TODO. These raw data samples are fed into the hardware first stage PFB, which carries out internal computation on TODO bit integer data, and generates 16-bit or 8-bit complex components (TODO: Is there any normalization at this stage? Why else are the actual numbers so high?). Due to each dipole sampled output being directly operated on without any beamforming, the 8-bit dynamic range has been found to be adequate for the AARTFAAC application. Once the data reaches the correlator, all computations are carried out with single precision floating points.

9. Conclusions

Acknowledgements. This work was funded by the ERC grant <num> awarded to Prof. Ralph Wijers, Universitiet Van Amsterdam. We thank The Netherlands Foundation for Radio Astronomy (ASTRON) for support provided in carrying out the commissioning observations.